ZS ASSOCIATES TECHNICAL ROUND

PROJECT: STATISTICAL DATA ANALYSIS ON UCI AIR QUALITY DATASET

Problem Statement

- The goal is to perform linear regression and time series analysis on the UCI Air quality dataset which contains 15 features and 9358 instances of hourly averaged responses from chemical sensors embedded in an Air Quality Chemical Multi sensor Device.

Data Description

Dataset : Air Quality dataset from UCI machine Learning Repository

Number of Instances: 9358 (hourly averaged responses of pollutants)

The dataset consists of following attributes: (Total :15)
Date, Time, CO(GT), PT08.S1(CO), NMHC(GT),C6H6(GT), PT08.S2(NHMC), NOx(GT), PT08.S3(NOx), NO2(GT), PT08.S4(NO2),PT08.S5(O3),T,RH and AH.

Dependent Variables: RH (Relative Humidity) and AH (Absolute Humidity)

Missing values are tagged by -200

Preprocessing

The missing values are tagged with -200.They are replaced with mean values of that particular day.

There are some Nan values left even after replacing with mean because there is no data available for whole day so we filled those Nan values with previous values.

There is one feature NHMC which has 90% of data composed of Nan. Replacing it with mean values is not a good idea as there are more Nan values.

As it doesn’t provide much information that attribute was dropped.

Preprocessing

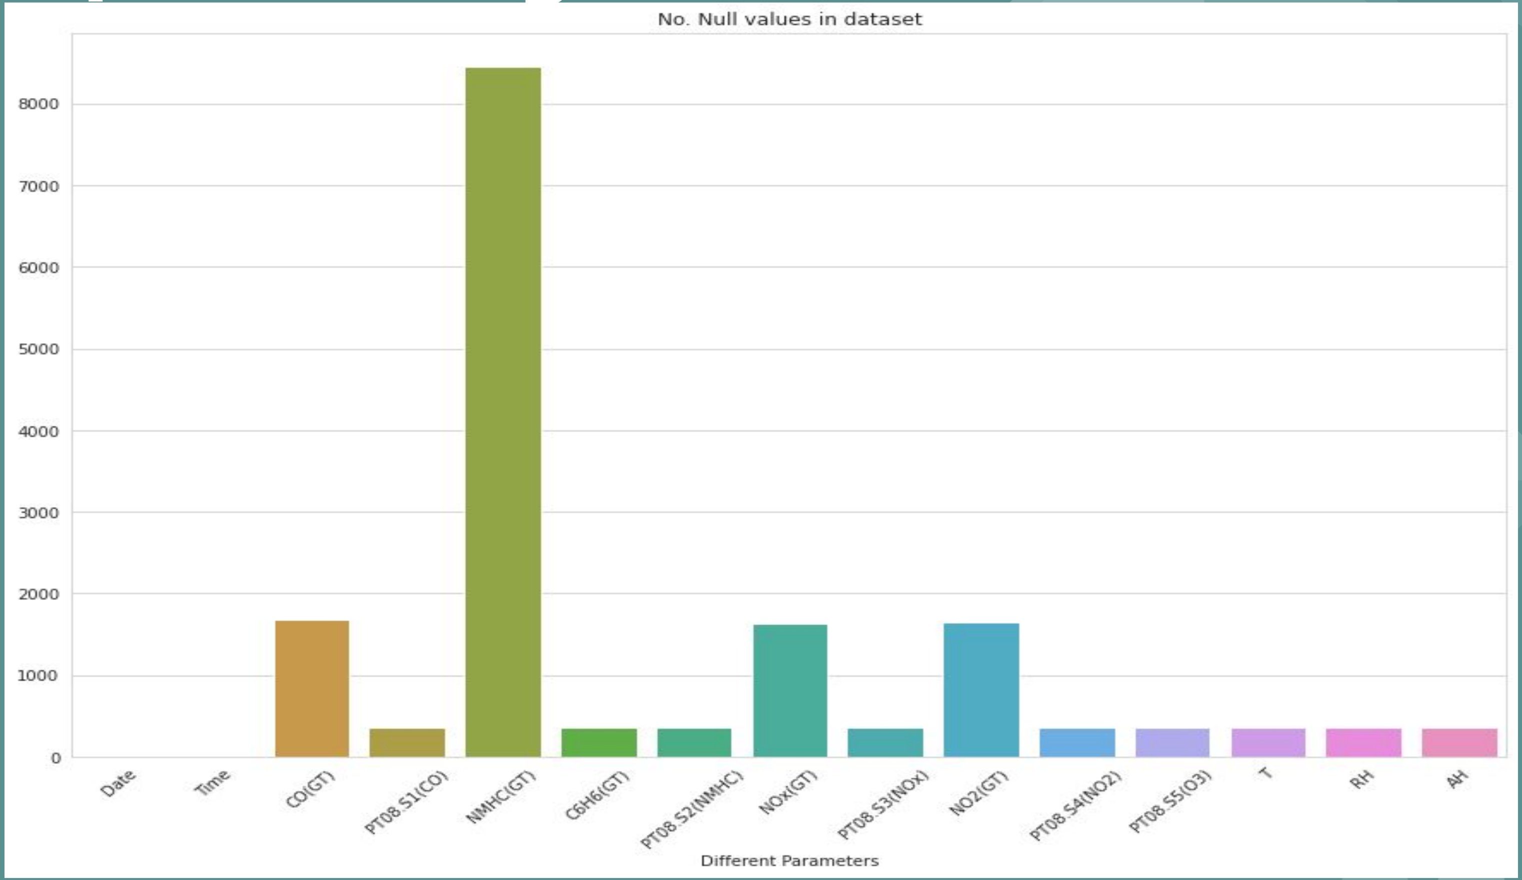

Data Analysis

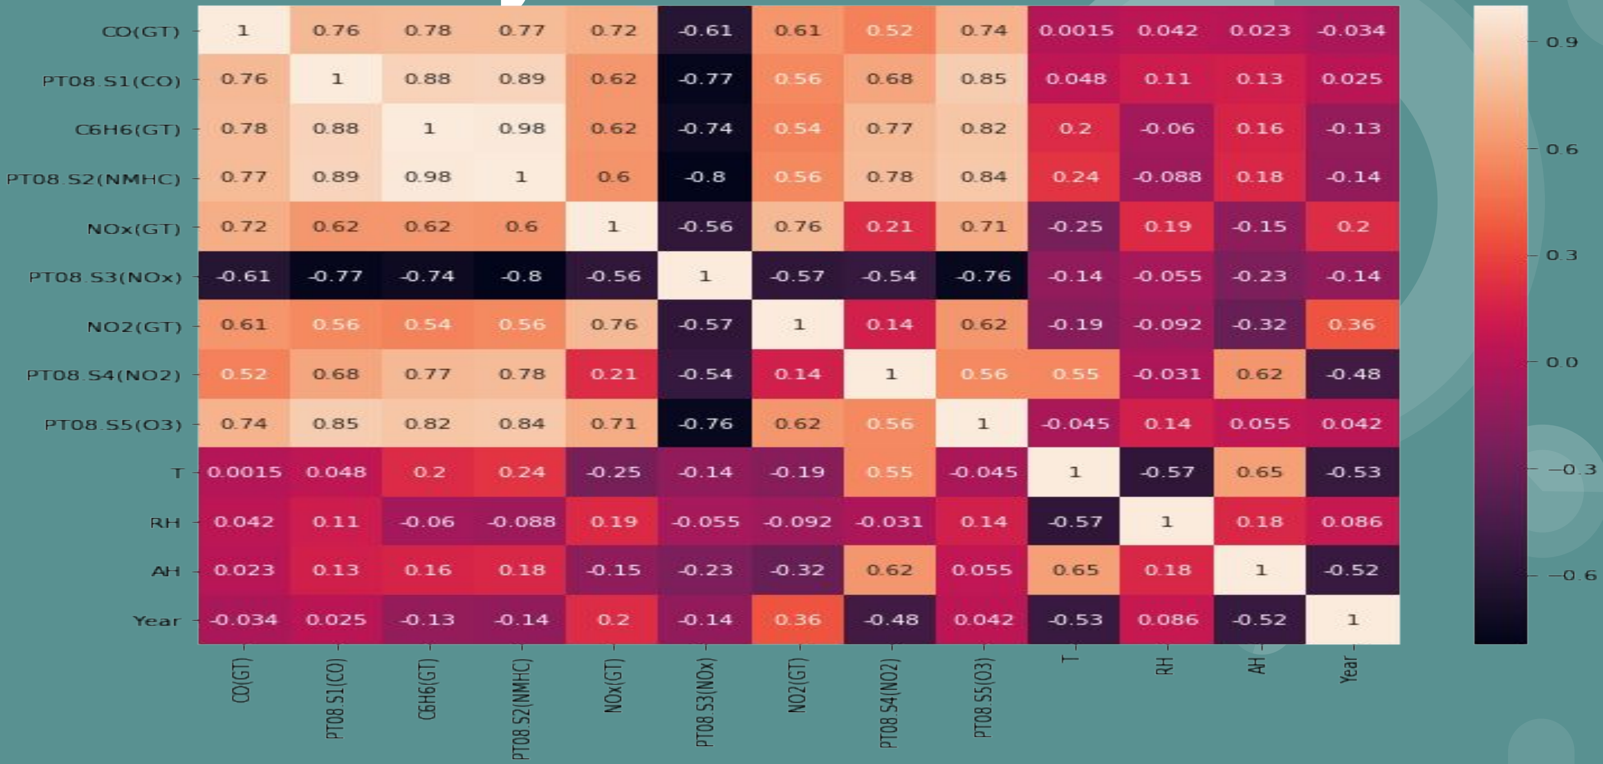

Above plot is the heatmap of correlation between all the variables.

DATA ANALYSIS

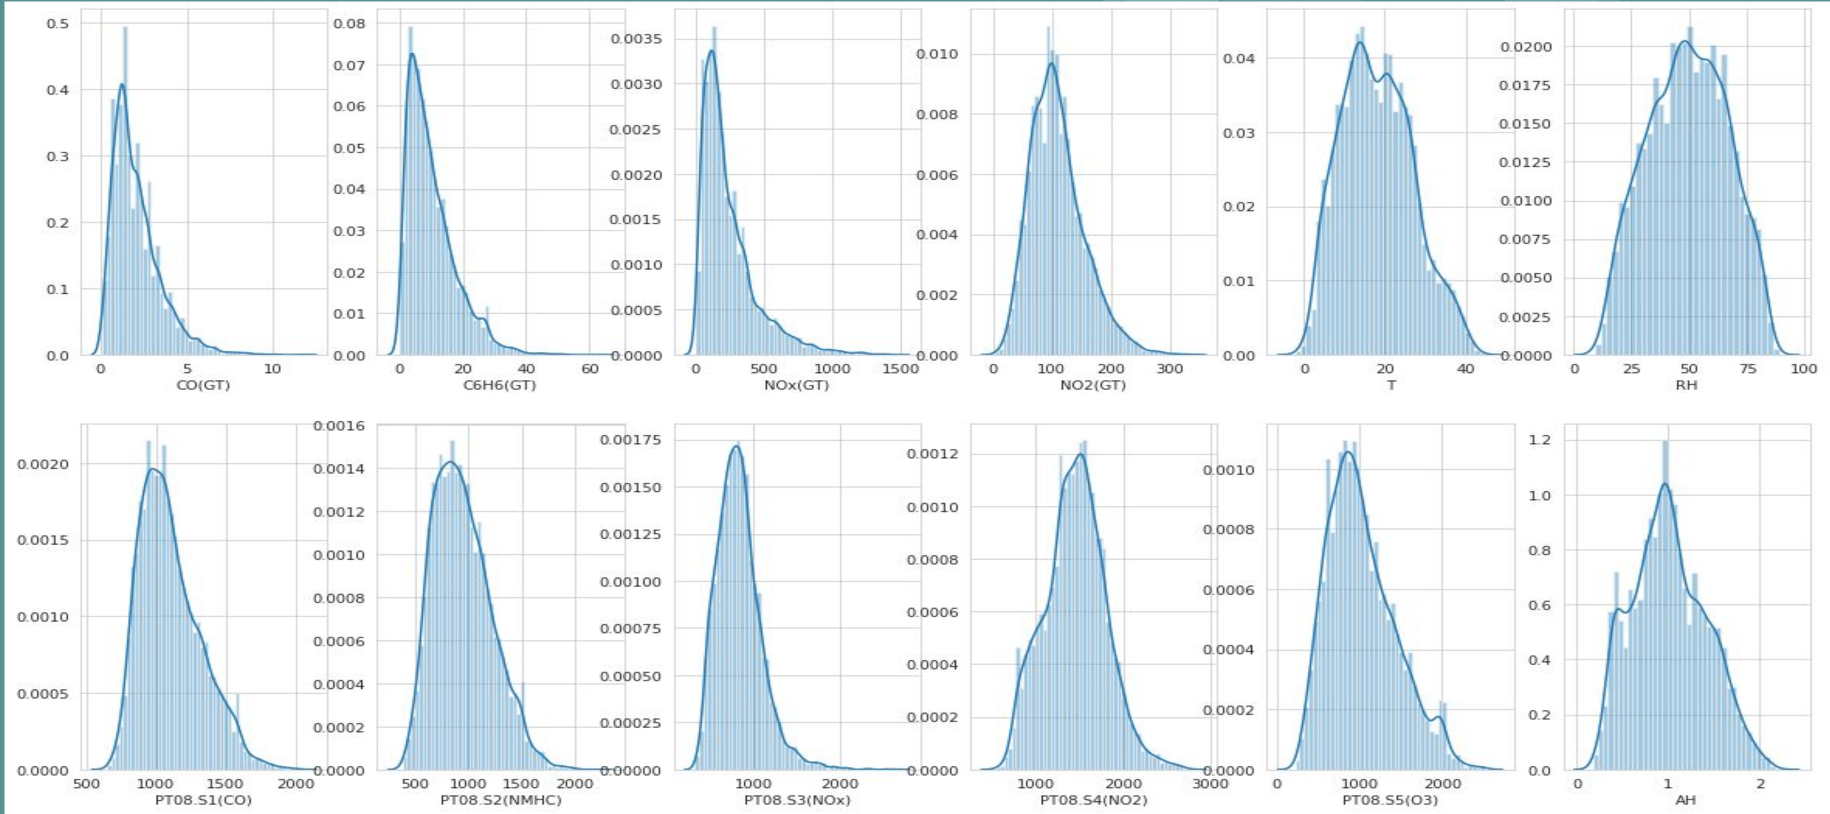

Data Analysis

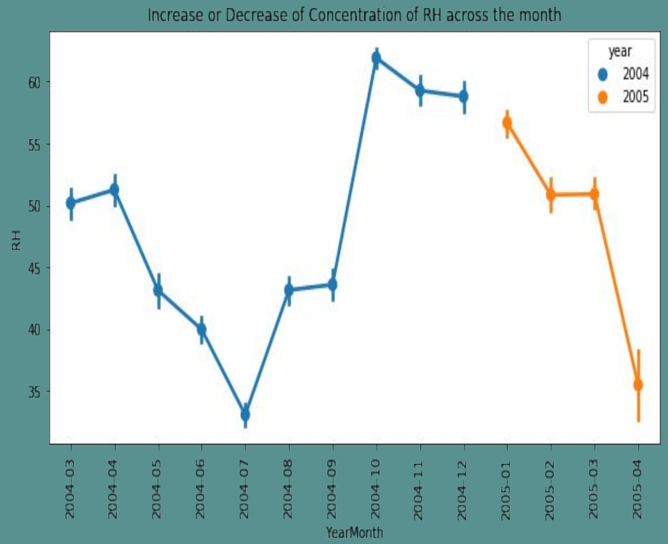

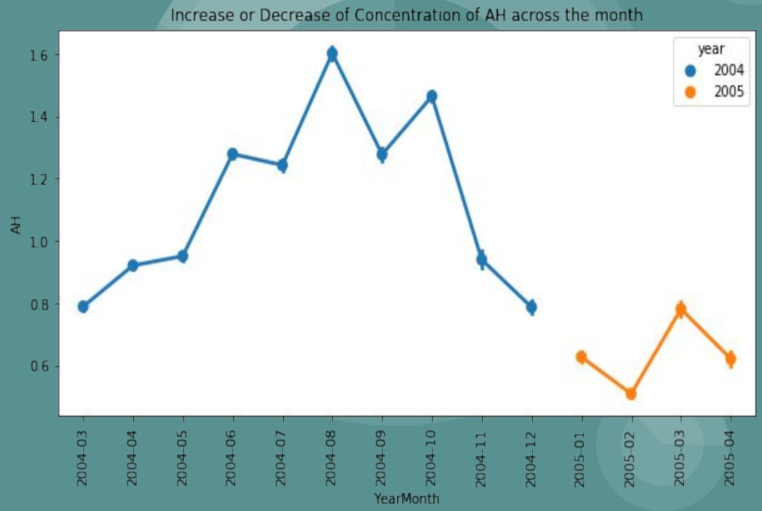

Above plots describes the variations of RH and AH along the year and increment or decrement of concentration across the months.

Assumptions of Multiple Linear Regression

- Multiple Linear Regression takes Several Assumptions
- It assumes Linear Relationship between response variable and independent variables. Scatter plots can show a linear or curvilinear relationship
- It assumes the residuals are normally distributed
- It assumes NO Multicollinearity which implies that independent variables should not be highly correlated with each other. It can be tested using Variance Inflation Factor(VIF) values
- It assumes Homoscedasticity.The assumption states that the variance of error terms is similar across different values of Independent variables.

Linearity Test:

Linear regression needs the relationship between the independent and dependent variables to be linear.

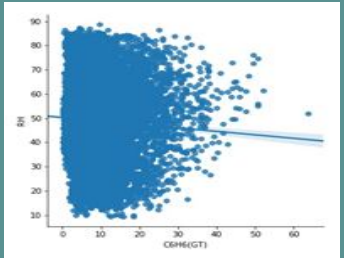

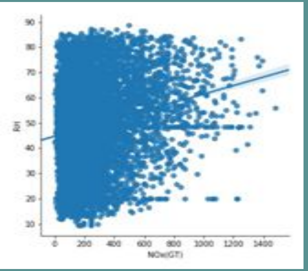

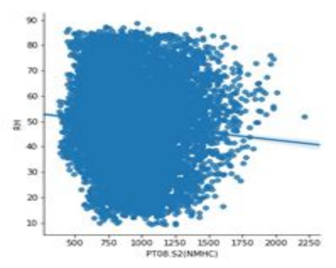

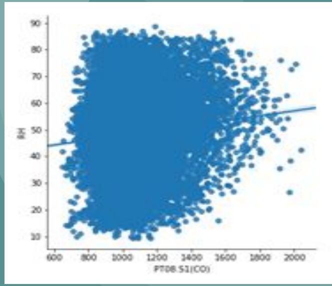

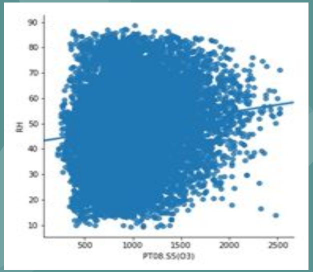

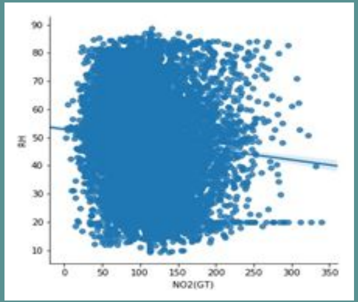

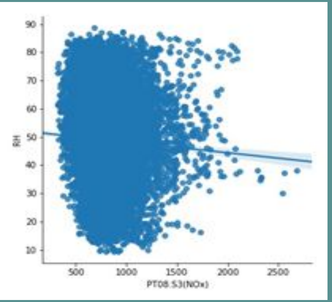

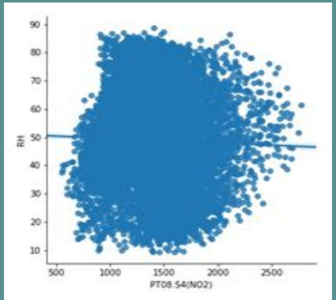

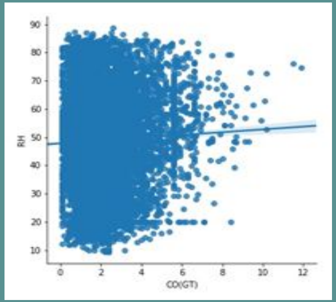

Normality Test:

The linear regression analysis needs all variables to be multivariate normal.
This assumption can best be checked with a histogram or a Q-Q-Plot graphically or normal test.

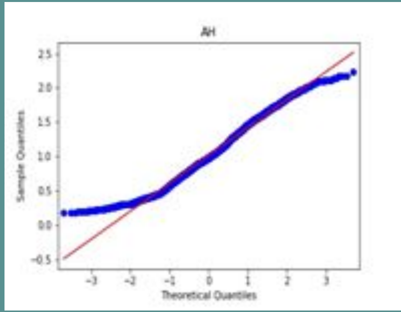

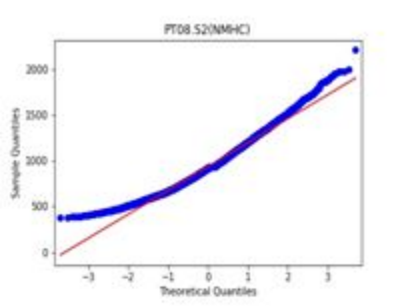

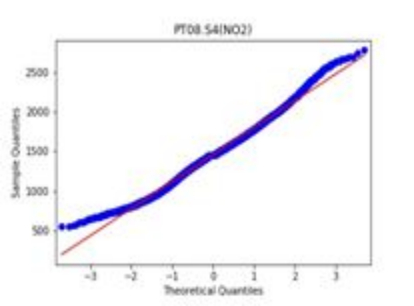

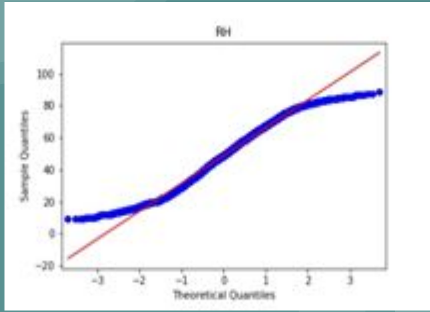

Some of the features are normal as it can be clearly seen from either of the graphs.
For remaining features, we can perform the stats.normaltest and find the corresponding p value.

NORMALITY TEST

- We can also use "normaltest" module from scipy package to test the normality.

Feature	P value
CO (GT)	0.0
T	8.877198092891439e-64
NO2 (GT)	8.382026215464602e-183
PT08.S4 (NO2)	3.1093274946139267e-18
PT08.S5 (O3)	4.489891015439565e-125
C6H6 (GT)	0.0
PT08.S2 (NMHC)	5.154651033295069e-100
PT08.S3 (NOx)	0.0
PT08.S1 (CO)	1.7886360129305134e-173
NOx (GT)	0.0
RH	7.022315598570015e-178
AH	2.1014760517340424e-74

From the table, we can see that all the p values are almost equal to 0 i.e., negligible value indicating that the data is normal.

Homoscedasticity:

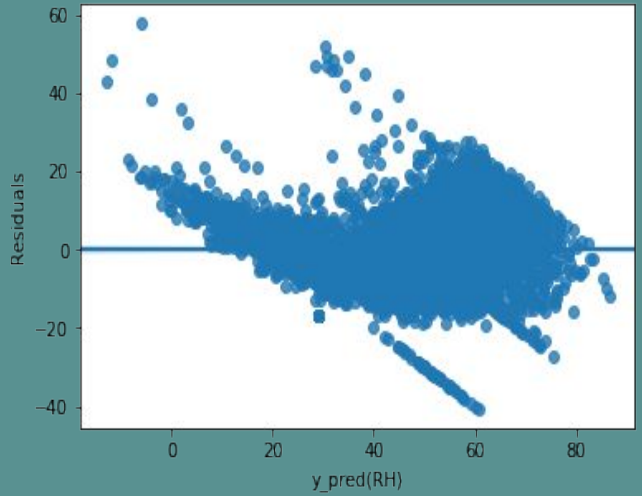

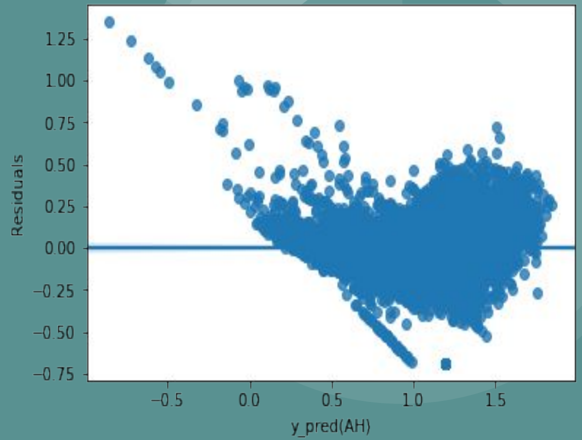

- The Above Plots conclude that the homoscedasticity assumption does not hold true.

Check for Stationarity

Dickey-Fuller Test:

The Dickey Fuller test is one of the most popular statistical tests. It can be used to determine the presence of unit root in the series, and hence help us understand if the series is stationary or not. The null and alternate hypothesis of this test are:

Null Hypothesis: The series has a unit root (value of a =1)

Alternate Hypothesis: The series has no unit root.

If we fail to reject the null hypothesis, we can say that the series is non-stationary. This means that the series can be linear or difference stationary

Check for Stationarity

Results of Dickey-Fuller Test:

Feature	P_value
RH	1.219023e-10
AH	0.000014
CO (GT)	5.412775e-16
T	0.019787
NO2 (GT)	7.786800e-13
PT08.S4 (NO2)	3.185933e-08
PT08.S5 (O3)	2.251934e-19
C6H6 (GT)	3.127256e-18
PT08.S2 (NMHC)	1.779690e-18
PT08.S3 (NOx)	5.035225e-19
PT08.S1 (CO)	8.914162e-17
NOx (GT)	2.985511e-11

From the above results the test statistic < critical value(alpha = 0.05), which implies that the series is stationary

Ordinary Least Squares(OLS):

Applying the ols regression from statsmodels.api, we get the following results.

For RH variable some of the results:

R-squared:	**0.736**
Durbin-Watson	1.993

For AH variable some of the results:

R-squared:	**0.798**
Durbin-Watson	**2.006**

Linear Regression

Linear regression is a statistical approach for modelling relationship between a dependent variable with a given set of independent variables.

After apply Linear Regression to our data by training the model using train data and predicting the test data set,we obtain the following results:

R2_score for the above model is 0.7587165188730305

FACTOR ANALYSIS

Factor Analysis (FA) is an exploratory data analysis method used to search influential underlying factors or latent variables from a set of observed variables. It extracts maximum common variance from all variables and puts them into a common score.

Adequacy Tests

- Bartlett's Test
- Kaiser-Meyer-Olkin Test

BARTLETT'S ADEQUACY TEST

Bartlett's test of sphericity checks whether or not the observed variables inter correlate at all.

If the test found statistically insignificant, you should not employ a factor analysis.

Result for Bartlett sphericity test:

We obtain p-value as 0.0

The test was statistically significant, indicating that the observed correlation matrix is not an identity matrix.Hence we can proceed with Factor Analysis.

kaiser-Meyer-Olkin Test

Kaiser-Meyer-Olkin (KMO) Test measures the suitability of data for factor analysis. If the Result obtained is greater than 0.6,we proceed with Factor Analysis.

Result:

0.8652245188839246

The overall KMO for our data is 0.8652, which is considerably good for the above test.Hence,we can proceed with Factor Analysis.

Next ,we perform the Factor Analysis for the data.

Factor Loadings of the Data

	Factor1	Factor2	Factor3	Factor4
NO2(GT)	0.245677	-0.102703	0.762700	0.212014
PT08.S2(NMHC)	0.867775	0.171821	0.421193	0.168409
PT08.S4(NO2)	0.770448	0.446725	0.043618	0.091084
PT08.S1(CO)	0.777560	-0.009076	0.426697	0.264350
NOx(GT)	0.290956	-0.154306	0.884593	0.034160
PT08.S5(O3)	0.661551	-0.065486	0.543830	0.284721
T	0.159564	0.969817	-0.169311	0.018212
CO(GT)	0.584358	0.010877	0.608930	0.043435
PT08.S3(NOx)	-0.572996	-0.120531	-0.446290	-0.525115
C6H6(GT)	0.893076	0.127648	0.424018	0.036974

Scree Plot

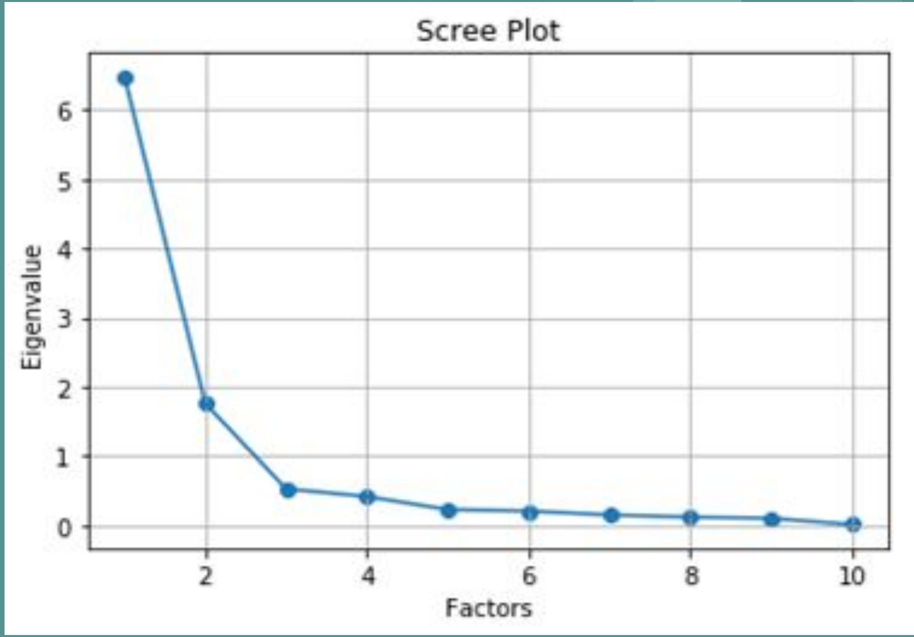

ARIMA

AH Variable:

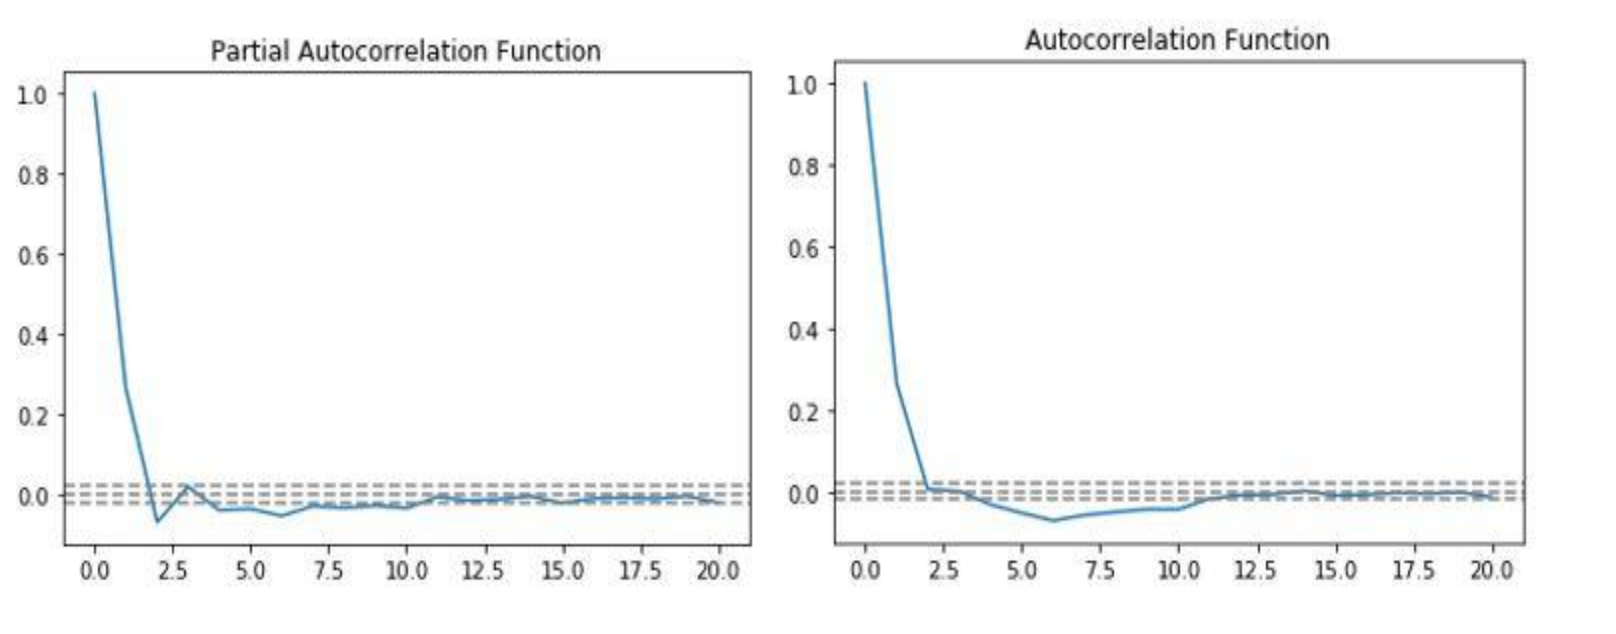

ARIMA

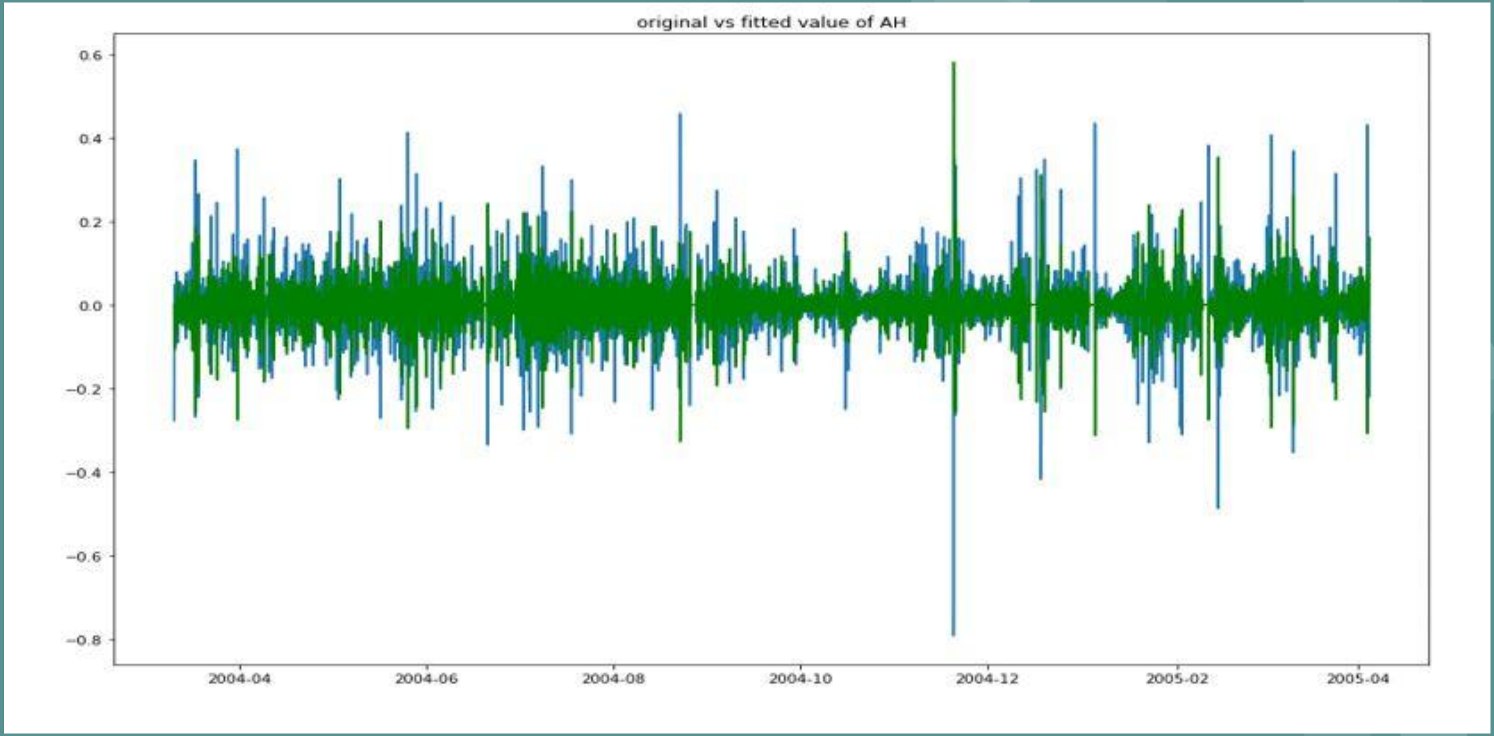
original vs fitted value of AH
0.6
0.4
0.2
0.0
−0.2
−0.4
−0.6
−0.8
2004-04
2004-06
2004-08
2004-10
2004-12
2005-02
2005-04

ARIMA

RH Variable

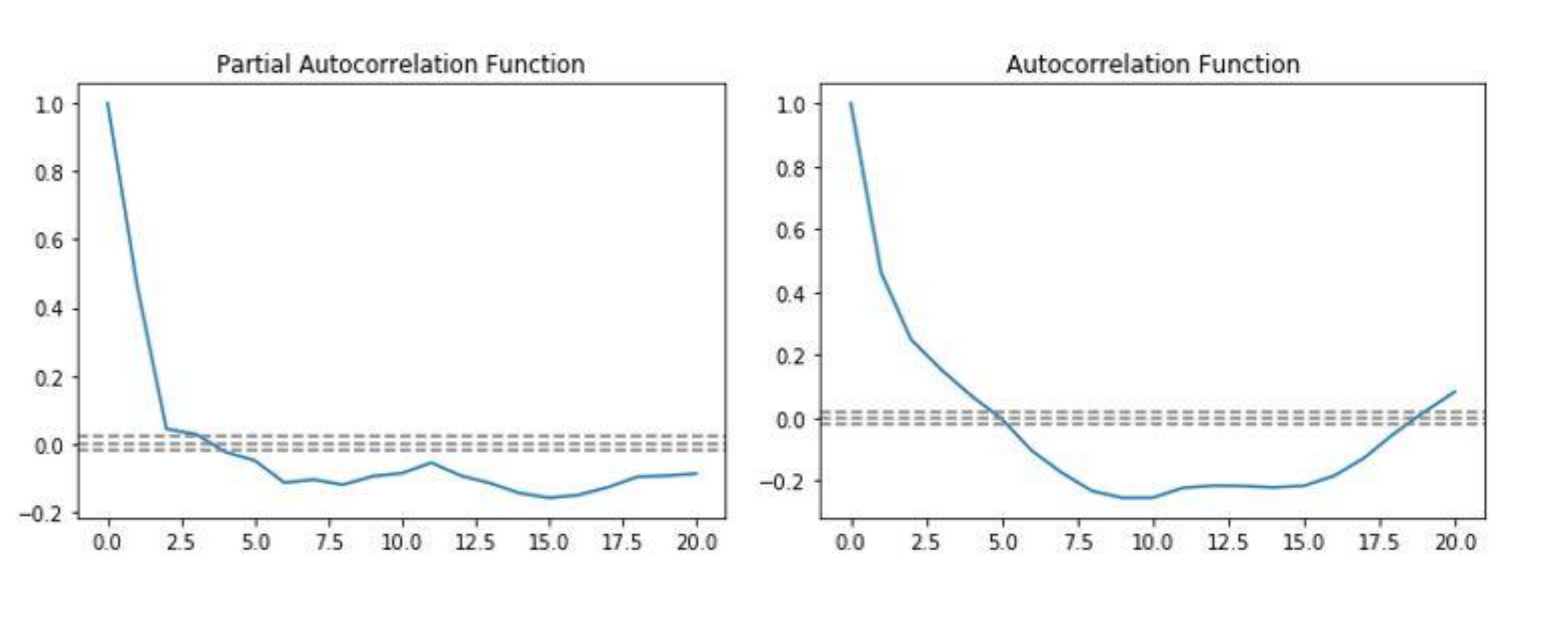

ARIMA

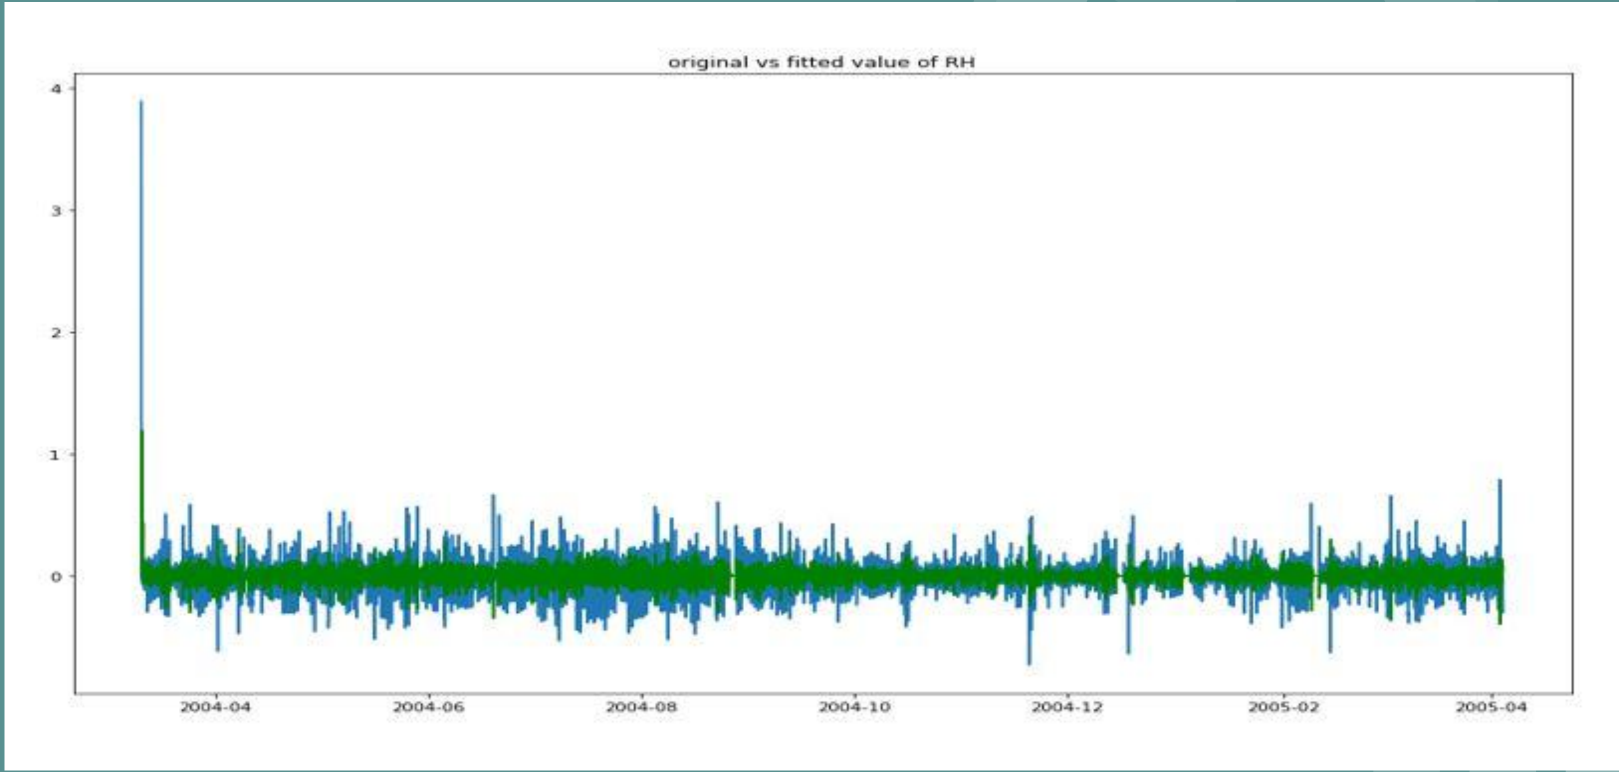
original vs fitted value of RH
4
3
2
1
0
2004-04
2004-06
2004-08
2004-10
2004-12
2005-02
2005-04